

# EECS 370 - Lecture 7

## Linking



# Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- Lecture 6 : Converting C to assembly – functions
- **Lecture 7 : Translation software; libraries, memory layout**



# Review

## Original C Code

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

No need to  
save r2/r3.  
Why?

## Additions for Caller-save

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    save r1 to stack  
    save r4 to stack  
    bar();  
    restore r1  
    restore r4  
    d = a+d;  
    return();  
}
```

Assume bar() will  
overwrite all registers

## Additions for Callee-save

```
void foo() {  
    int a,b,c,d;  
    save r1  
    save r2  
    save r3  
    save r4  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    bar();  
    d = a+d;  
    restore r1  
    restore r2  
    restore r3  
    restore r4  
    return();  
}
```

bar() will save a,b, but  
now foo() must save  
main's variables

# Question 1: Caller-save

```
void main() {  
    int a,b,c,d;  
    c = 5; d = 6;  
    a = 2; b = 3;  
    [4 STUR]  
    foo();  
    [4 LDUR]  
    d = a+b+c+d;  
}
```

```
void foo() {  
    int e,f;  
  
    e = 2; f = 3;  
    [2 STUR]  
    bar();  
    [2 LDUR]  
    e = e + f;  
}
```

```
void bar() {  
    int g,h,i,j;  
    g = 0; h = 1;  
    i = 2; j = 3;  
    [3 STUR]  
    final();  
    [3 LDUR]  
    j = g+h+i;  
}
```

```
void final() {  
    int y,z;  
  
    y = 2; z = 3;  
  
    z = y+z;  
}
```

Total: 9 STUR / 9 LDUR

## Question 2: Callee-save

```
void main() {  
    int a,b,c,d;  
  
    c = 5; d = 6;  
    a = 2; b = 3;  
    foo();  
    d = a+b+c+d;  
  
}
```

```
void foo() {  
    [2 STUR]  
    int e,f;  
  
    e = 2; f = 3;  
    bar();  
    e = e + f;  
  
    [2 LDUR]  
}
```

```
void bar() {  
    [4 STUR]  
    int g,h,i,j;  
    g = 0; h = 1;  
    i = 2; j = 3;  
    final();  
    j = g+h+i;  
  
    [4 LDUR]  
}
```

```
void final() {  
    [2 STUR]  
    int y,z;  
  
    y = 2; z = 3;  
  
    z = y+z;  
  
    [2 LDUR]  
}
```

Total: 8 STUR / 8 LDUR

# Is one better?

- **Caller-save** works best when we don't have many live values across function call
- **Callee-save** works best when we don't use many registers overall
- We probably see functions of both kinds across an entire program
- Solution:
  - Use both!
  - E.g. if we have 6 registers, use some (say r0-r2) as **caller-save** and others (say r3-r5) as **callee-save**
  - Now each function can optimize for each situation to reduce saving/restoring
  - Not discussed further in this class

# LEGv8 ABI- Application Binary Interface

- The ABI is an agreement about how to use the various registers
- Not enforced by hardware, just a convention by programmers / compilers
- If you want your code to work with other functions / libraries, **follow these**
- Some register conventions in ARMv8
  - X30 is the **link register** – used to hold return address
  - X28 is **stack pointer** – holds address of top of stack
  - X19-X27 are **callee-saved** – function must save these before writing to them
  - X0-15 are **caller-saved** –function must save live values before call
  - X0-X7 used for **arguments** (memory used if more space is needed)
  - X0 used for **return value**

# Caller/Callee

- Still not clicking?
- Don't worry, this is a tricky concept for students to get
- Check out supplemental video
  - <https://www.youtube.com/watch?v=SMH5uL3HiiU>
  - Don't worry about mixed caller/callee save case at the end
- Come to office hours to go over examples



# Today we'll finish up software

- Introduce linkers and loaders
  - Basic relationship of compiler, assembler, linker and loader.
  - Object files
    - Symbol tables and relocation tables

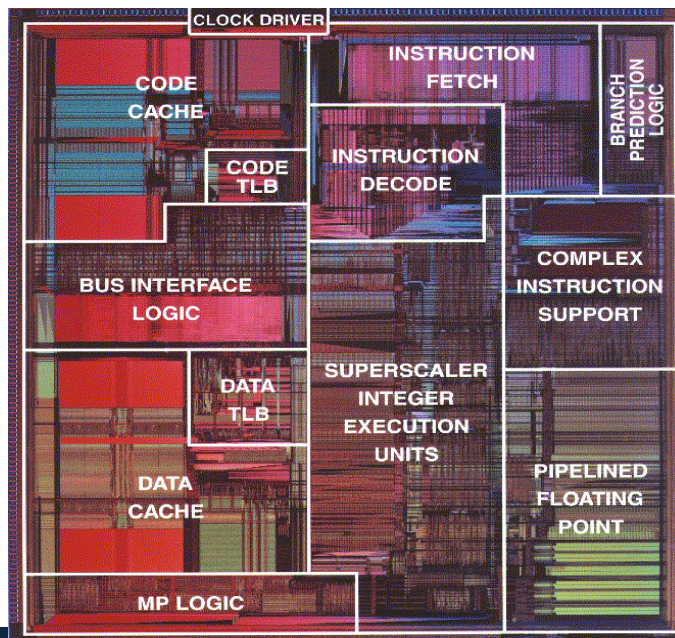
# Source Code to Execution

- In project 1a, our view is this:

Assembly

Assembler

Poll: Why do we write programs in multiple files?



Executable

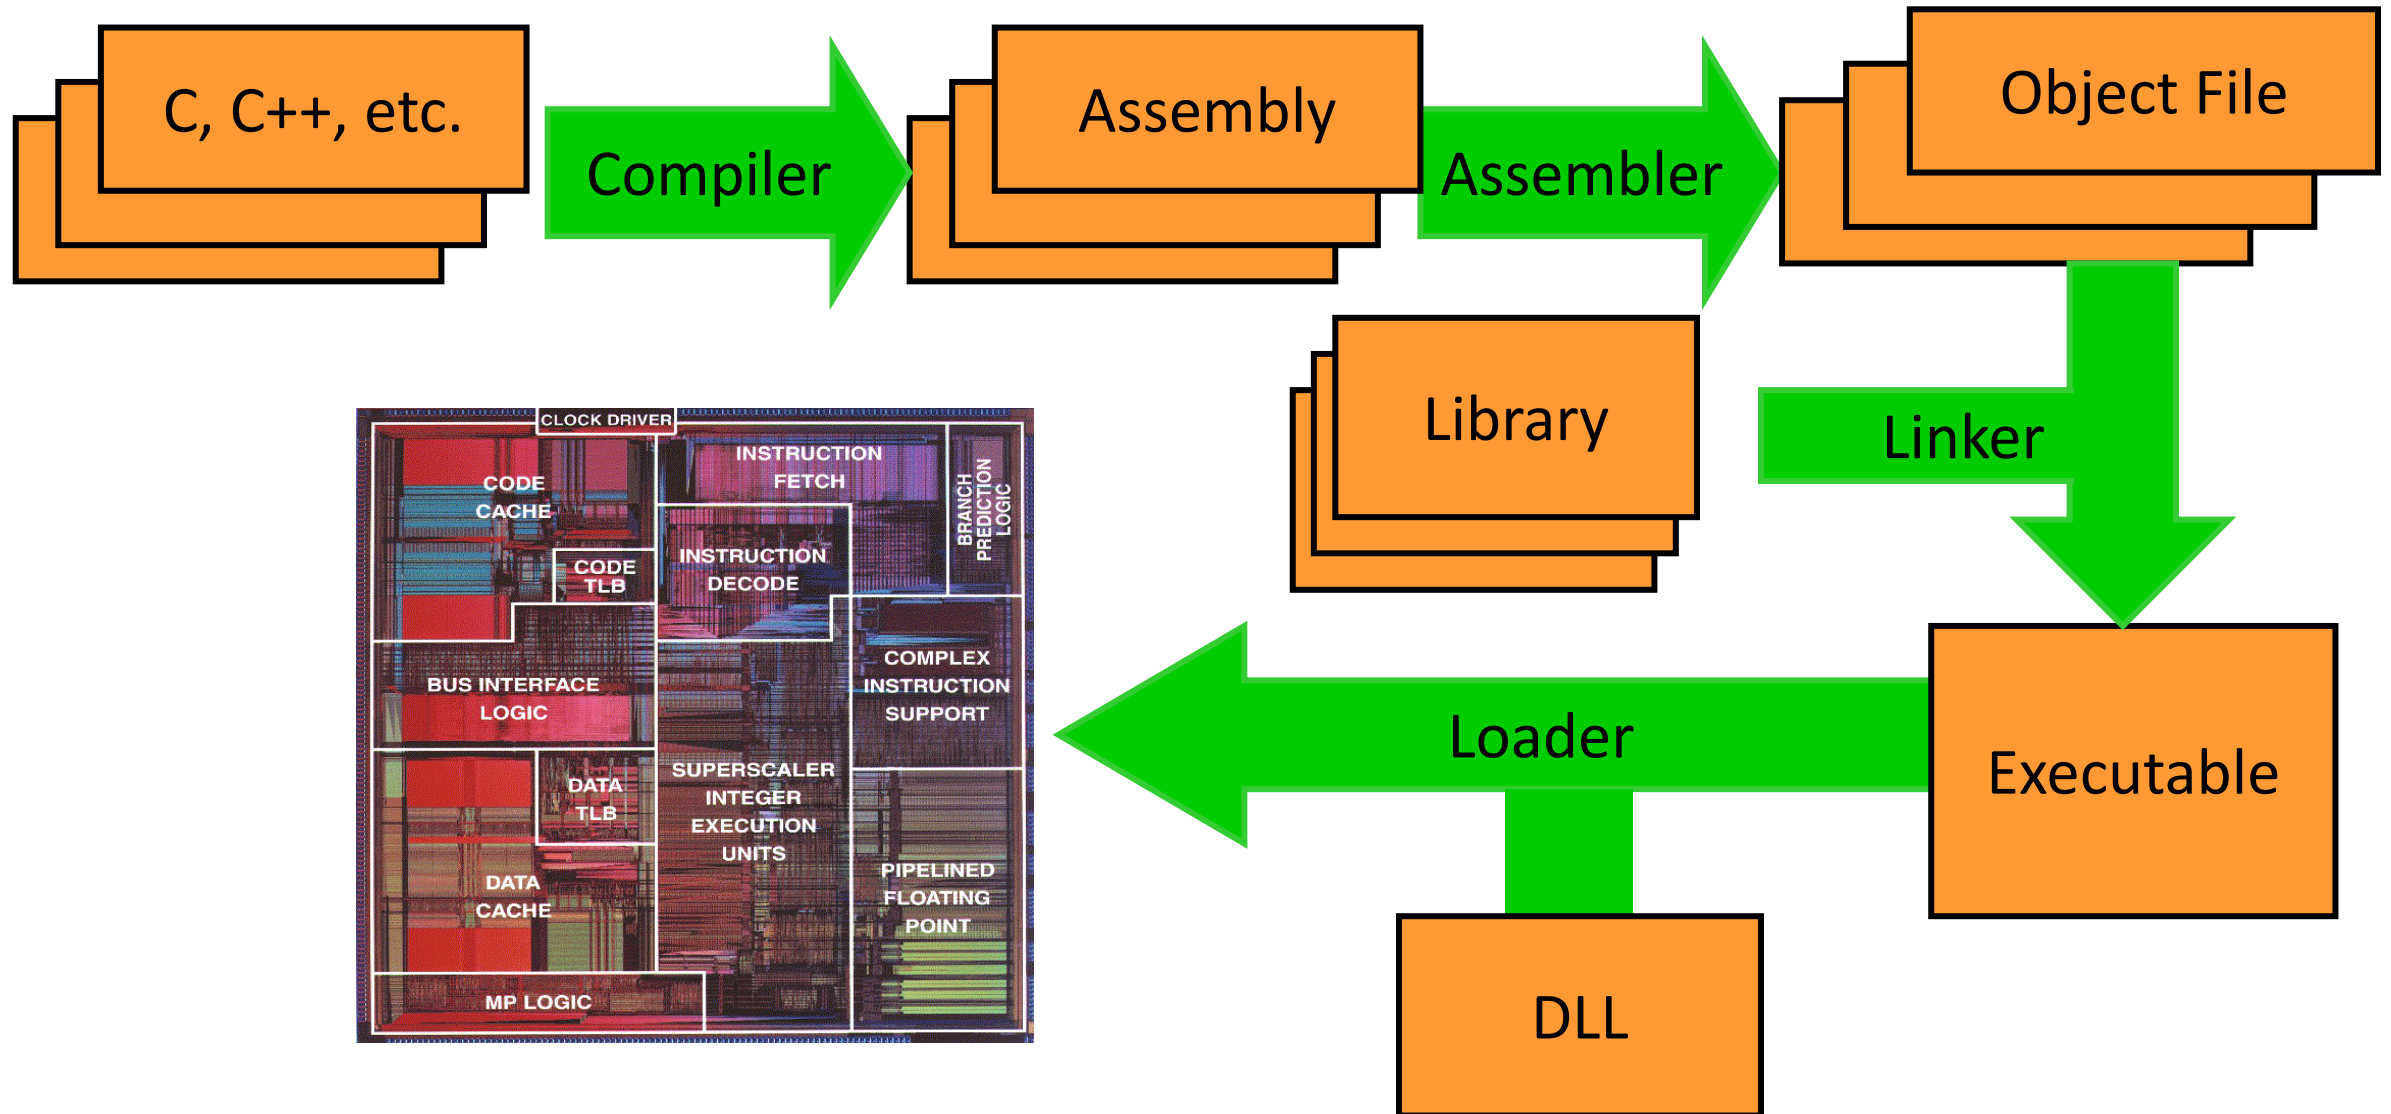
Not very accurate... why?  
Because in reality, we have  
multiple files

# Multi-file programs

- In practice, programs are made from thousands or millions of lines of code
  - Use pre-existing libraries like stdlib
- If we change one line, do we need to recompile the whole thing?
  - No! If we compile each file into a separate **object file**, then we only need to recompile that one file and **link** it to the other, unchanged object files

# Source Code to Execution

What do object files look like?



# What do object files look like?

```
extern int X;  
extern void foo();  
int Y;  
  
void main() {  
    Y = X + 1;  
    foo();  
}
```

"extern" means  
defined in another  
file

```
extern int Y;  
int X;  
  
void foo() {  
    Y *= 2;  
}
```

Compile →

```
.main:  
LDUR    X1, [XZR, X]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, Y]  
BL      foo  
HALT
```

Compile →

```
.foo:  
LDUR    X1, [XZR, Y]  
LSL     X9, X1, #1  
STUR    X9, [XZR, Y]  
BR      X30
```

Uh-oh!  
Don't know  
address of X, Y,  
or foo!

# Linking

```
.main:  
LDUR    X1, [XZR, X]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, Y]  
BL      foo  
HALT
```

```
.foo:  
LDUR    X1, [XZR, Y]  
LSL     X9, X1, #1  
STUR    X9, [XZR, Y]  
BR      X30
```

Assemble

???

Assemble

???

What needs to go  
in this intermediate  
"object file"?

LINK

LINK

NOTE: this will  
actually be in  
machine code, not  
assembly

```
LDUR    X1, [XZR, #40]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, #36]  
BL      #2  
HALT  
LDUR    X1, [XZR, #36]  
LSL     X9, X1, #1  
STUR    X9, [XZR, #36]  
BR      X30  
// Addr #36 starts here
```

# Linking

.main:

```
LDUR    X1, [XZR, X]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, Y]  
BL      foo  
HALT
```

Assemble

???

LINK

```
LDUR    X1, [XZR, #40]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, #36]  
BL      #2  
HALT  
LDUR    X1, [XZR, #36]  
LSL     X9, X1, #1  
STUR    X9, [XZR, #36]  
BR      X30
```

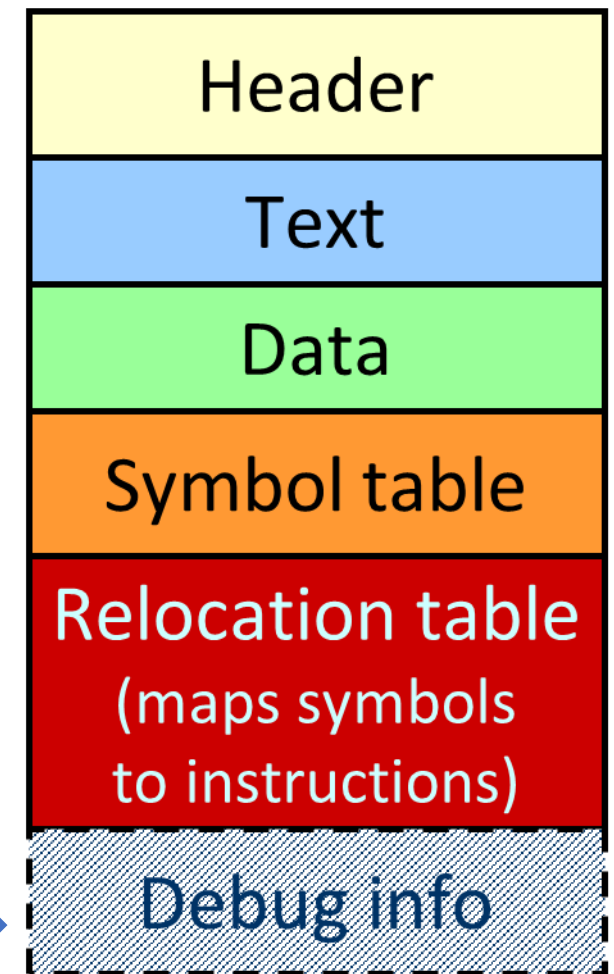
We need:

- the assembled machine code:
- list of instructions that need to be updated once addresses are resolved
- list of symbols for cross-ref

# What do object files look like?

- Since we can't make executable, we make an object file
- Basically, includes the machine code that will go in the executable
  - Plus extra information on what we need to modify once we stitch all the other object files together
- Looks like this ->

## Object code format



We won't discuss "Debug" much. Gets included when you compile with "-g" in gcc



# Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    Y = G + 1;
    B();
}
```

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

<b>Header</b>	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
<b>Text</b>	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
<b>Data</b>	0	X	3
<b>Symbol table</b>	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
<b>Reloc table</b>	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

# Assembly → Object file - example

```
extern int X;
extern void B();
int X = 3;
main() {
    Y = G + 1;
    B();
}
```

**Header:**  
keeps track of  
size of each  
section

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

<b>Header</b>	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
<b>Text</b>	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
<b>Data</b>	0	X	3
<b>Symbol table</b>	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
<b>Reloc table</b>	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

# Assembly → Object file - example

```
extern int G;  
extern void B();  
int X = 3;  
main() {  
    Y = G +  
    B();  
}
```

**Text:**  
machine code

```
LDUR    X1, [XZR, G]  
ADDI    X9, X1, #1  
BL      B
```

<b>Header</b>	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
<b>Text</b>	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
<b>Data</b>	0	X	3
<b>Symbol table</b>	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
<b>Reloc table</b>	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

Simplifying Assumption for EECS370

All globals and static locals (initialized or not) go in the data segment

Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    Y = G + 1;
    B();
}
```

**Data:**  
initialized globals  
and static locals

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

Header	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
Text	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
Data	0	X	3
Symbol table	Label	Address	
	X	0	
	B	-	
	main	0	
Reloc table	G	-	
	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B



# Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    Y = G + 1;
    B();
}
```

LDUR  
ADDI  
BL

**Symbol table:**  
Lists all labels  
visible outside this file  
(i.e. function names  
and global variables)

<b>Header</b>	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
<b>Text</b>	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
<b>Data</b>	0	X	3
<b>Symbol table</b>	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
<b>Reloc table</b>	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

# Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    Y = G + 1;
    B();
}
```

LDUR X1, [XZR, G]

**Relocation Table:**  
list of instructions and data  
words that must be updated  
if things are moved in  
memory

<b>Header</b>	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
<b>Text</b>	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
<b>Data</b>	0	X	3
<b>Symbol table</b>	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
<b>Reloc table</b>	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

# Class Problem 1

**Poll: Which symbols will be put in the symbol table?** (i.e. which "things" should be visible to all files?)

```
file1.c
extern void bar(int);
extern char c[];
int a;
int foo (int x) {
    int b;
    a = c[3] + 1;
    bar(x);
    b = 27;
}
```

file 1 – symbol table

<b>sym</b>	<b>loc</b>
a	data
foo	text
c	-
bar	-

```
file2.c
extern int a;
char c[100];
void bar (int y) {
    char e[100];
    a = y;
    c[20] = e[7];
}
```

file 2 – symbol table

<b>sym</b>	<b>loc</b>
c	data
bar	text
a	-

**Poll: Which lines / instructions are in the relocation table?** (i.e. which "things" need to be updated after linking?)

## Class Problem 2

file1.c

```
1 extern void bar(int);
2 extern char c[];
3 int a;
4 int foo (int x) {
5     int b;
6     a = c[3] + 1;
7     bar(x);
8     b = 27;
9 }
```

file 1 - relocation table

line	type	dep
6	ldur	c
6	stur	a
7	bl	bar

file2.c

```
1 extern int a;
2 char c[100];
3 void bar (int y) {
4     char e[100];
5     a = y;
6     c[20] = e[7];
7 }
```

file 2 - relocation table

line	type	dep
5	stur	a
6	stur	c

Note: in a real relocation table, the "line" would really be the address in "text" section of the instruction we need to update.



# Linker

- Stitches independently created object files into a single executable file (i.e., a.out)
  - Step 1: Take text segment from each .o file and put them together.
  - Step 2: Take data segment from each .o file, put them together, and concatenate this onto end of text segments.
- What about libraries?
  - Libraries are just special object files.
  - You create new libraries by making lots of object files (for the components of the library) and combining them (see ar and ranlib on Unix machines).
- Step 3: Resolve cross-file references to labels
  - Make sure there are no undefined labels

# Linker - Continued

- Determine the memory locations the code and data of each file will occupy
  - Each function could be assembled on its own
  - Thus, the relative placement of code/data is not known up to this point
  - **Must relocate absolute references to reflect placement by the linker**
    - PC-Relative Addressing (beq, bne): never relocate
    - Absolute Address (mov 27, #X): always relocate
    - External Reference (usually bl): always relocate
    - Data Reference (often movz/movk): always relocate
- Executable file contains no relocation info or symbol table  
these just used by assembler/linker

# Loader

- Executable file is sitting on the disk
- Puts the executable file code image into memory and asks the operating system to schedule it as a new process
  - Creates new address space for program large enough to hold text and data segments, along with a stack segment
  - Copies instructions and data from executable file into the new address space
  - Initializes registers (PC and SP most important)
- Take operating systems class (EECS 482) to learn more!

# Summary

- Compiler converts a single source code file into a single assembly language file
- Assembler handles directives (.fill), converts what it can to machine language, and creates a checklist for the linker (relocation table). This changes each .s file into a .o file
- Assembler does 2 passes to resolve addresses, handling internal forward references
- Linker combines several .o files and resolves absolute addresses
- Linker enables separate compilation: Thus unchanged files, including libraries need not be recompiled.
- Linker resolves remaining addresses.
- Loader loads executable into memory and begins execution

# Floating Point Arithmetic

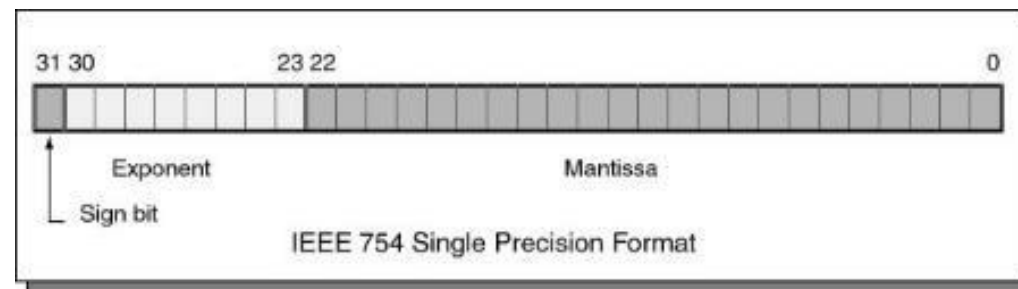


# Why floating point

- Have to represent real numbers somehow
- Rational numbers
  - Ok, but can be cumbersome to work with
- Fixed point
  - Do everything in thousandths (or millionths, etc.)
  - Not always easy to pick the right units
  - Different scaling factors for different stages of computation
- **Scientific notation: this is good!**
  - Exponential notation allows HUGE dynamic range
  - Constant (approximately) relative precision across the whole range

# IEEE Floating point format (single precision)

- Sign bit: (0 is positive, 1 is negative)
- Significand: (also called the *mantissa*; stores the 23 most significant bits after the decimal point)
- Exponent: used biased base 127 encoding
  - Add 127 to the value of the exponent to encode:
  - -127 → 00000000      1 → 10000000
  - -126 → 00000001      2 → 10000001
  - ...
  - 0 → 01111111      128 → 11111111
- How do you represent zero ? Special convention:
  - Exponent: -127 (all zeroes ), Significand 0 (all zeroes), Sign + or -



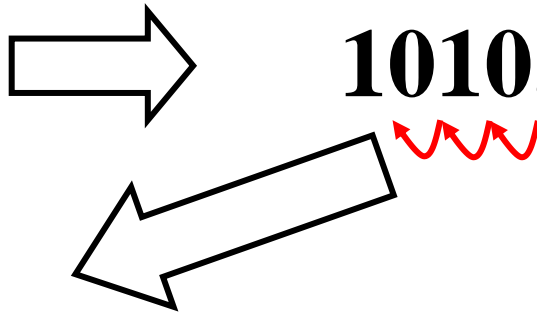
# Floating Point Representation

$$10.625_{10} \Rightarrow 1010.101_2$$

- ❑ Step 1: convert from decimal to binary
  - 1<sup>st</sup> bit after "binary" point represents 0.5 (i.e.  $2^{-1}$ )
  - 2<sup>nd</sup> bit represents 0.25 (i.e.  $2^{-2}$ )
  - etc.



# Floating Point Representation

$$10.625_{10} \quad \Rightarrow \quad 1010.101_2$$


$$1.010101 \times 2^3$$

- Step 2: normalize number by shifting binary point until you get  $1.XXX \times 2^Y$

# Floating Point Representation

$$10.625_{10} \longrightarrow 1010.101_2$$

$$1.010101 \times 2^3$$

This must be a 1!  
So don't store it.



- Step 3: store relevant numbers in proper location (ignoring initial 1 of significand)

# Floating Point Representation

$$10.625_{10} \longrightarrow 1010.101_2$$

$$1.010101 \times 2^3$$

This must be a 1!  
So don't store it.



$$10.625_{10} = 0 \ 10000010 \ 010101000000000000000000$$

# Next Time

- Wrap up Floating Point
- And... hardware time!